



# Week 3 – Data Manipulation Language



## W3 – Housekeeping

- ➊ Using own computer (in Labs or Campus) - please log in to partch
  - Instruction in Lab 1
- ➋ Thanks Patrik for providing More SQL exercises
  - Find in Wee3 learning materials on Wattle
- ➌ Drop-in session on Tuesday 11am-12pm at CSIT N19.



## Week 3

- Quiz feedback
- Insert, delete, update
- Select: simple SQL queries
- Select: advanced SQL queries



## Quiz-Q5

### Q5 Version 1 (latest)

Question 1

Not yet  
answered

Marked out of  
0.10

Given the following two relations, and the SQL query

```
SELECT *  
FROM ANIMAL INNER JOIN COLOR  
ON ANIMAL.B = COLOR.E;
```

ANIMAL		
A	B	C
1	white	cat
2	brown	rabbit
3	white	bird
4	red	bird

COLOR	
D	E
1	brown
2	white
3	blue

How many rows will be in the resulting table of the above SQL query?

- a. 4
- b. None of the above
- c. 5
- d. 3



## Quiz-Q6 statistics

### Analysis of responses

Part of question	Response	Partial credit	Count	Frequency
102655908	SQL (Part 2): Insert, Update, Delete	0.00%	22	5.68%
102655909	SQL (Part 3): Simple SQL Queries	0.00%	30	7.75%
102655910	SQL (Part 4): Advanced SQL Queries	0.00%	236	60.98%
102655911	None of the above	100.00%	132	34.11%



## Week 3

- Quiz feedback
- Insert, delete, update
- Select: simple SQL queries
- Select: advanced SQL queries



## Insert, Update, Delete

- The **INSERT** statement is used to add tuples into a relation.

```
INSERT INTO table_name
    [(attribute_name, . . . , attribute_name)]
VALUES (value, . . . , value), . . . , (value, . . . , value);
```

- The **UPDATE** statement is used to modify attribute values of one or more selected tuples.

```
UPDATE table_name
    SET attribute_name = value, . . . , attribute_name = value
    [WHERE selection_condition];
```

- The **DELETE** statement is used to remove tuples from a relation.

```
DELETE FROM table_name
    [WHERE selection_condition];
```



## Insert – Example

```
CREATE TABLE STUDENT(StudentID INT PRIMARY KEY, Name VARCHAR(50),  
DoB DATE, Email VARCHAR(100));
```

- Will the following Insert statements work?

- ```
INSERT INTO STUDENT  
VALUES (456, 'Tom', '25/01/1988', 'tom@gmail.com');
```

**Yes.**

- ```
INSERT INTO STUDENT(StudentID)  
VALUES (459);
```

**Yes.** The values for Name, DoB and Email will be NULL.

- ```
INSERT INTO STUDENT(Name, DoB, Email)  
VALUES ('John', '15/11/1998', 'john@gmail.com');
```

**No.** The primary key value cannot be NULL.



## Update – Example

| STUDENT   |       |            |                  |
|-----------|-------|------------|------------------|
| StudentID | Name  | DoB        | Email            |
| 456       | Tom   | 25/01/1988 | tom@gmail.com    |
| 458       | Peter | 23/05/1993 | peter@gmail.com  |
| 459       | Fran  | 11/09/1987 | frankk@gmail.com |

- What is the resulting table after executing the following statement?

```
UPDATE STUDENT SET Name='Tom Lee', Email='tom.lee@yahoo.com'  
WHERE StudentID=456;
```

| STUDENT   |         |            |                   |
|-----------|---------|------------|-------------------|
| StudentID | Name    | DoB        | Email             |
| 456       | Tom Lee | 25/01/1988 | tom.lee@yahoo.com |
| 458       | Peter   | 23/05/1993 | peter@gmail.com   |
| 459       | Fran    | 11/09/1987 | frankk@gmail.com  |



## Delete – Example

| STUDENT   |       |            |                  |
|-----------|-------|------------|------------------|
| StudentID | Name  | DoB        | Email            |
| 456       | Tom   | 25/01/1988 | tom@gmail.com    |
| 458       | Peter | 23/05/1993 | peter@gmail.com  |
| 459       | Fran  | 11/09/1987 | frankk@gmail.com |

- What is the resulting table after executing the following statement?

```
DELETE FROM STUDENT WHERE StudentID=456;
```

| STUDENT   |       |            |                  |
|-----------|-------|------------|------------------|
| StudentID | Name  | DoB        | Email            |
| 458       | Peter | 23/05/1993 | peter@gmail.com  |
| 459       | Fran  | 11/09/1987 | frankk@gmail.com |

```
DELETE FROM STUDENT;
```

| STUDENT   |      |     |       |
|-----------|------|-----|-------|
| StudentID | Name | DoB | Email |
|           |      |     |       |

```
DROP TABLE STUDENT;
```

The Table STUDENT is deleted.

What is the difference?: DELETE/ DROP, UPDATE/ALTER



## Update and Delete - Referential Actions

| ENROL     |          |          |        |            |
|-----------|----------|----------|--------|------------|
| StudentID | CourseNo | Semester | Status | EnrolDate  |
| 456       | COMP1130 | 2016 S1  | active | 25/02/2016 |
| 458       | COMP1130 | 2016 S1  | active | 25/02/2016 |
| 456       | COMP2400 | 2016 S2  | active | 09/03/2016 |

| STUDENT   |       |            |                   |
|-----------|-------|------------|-------------------|
| StudentID | Name  | DoB        | Email             |
| 456       | Tom   | 25/01/1988 | tom@gmail.com     |
| 458       | Peter | 20/02/1991 | peter@hotmail.com |

- Referential actions specify what happens in case of deleting or updating referenced tuples. SQL offers the following possibilities:
  - NO ACTION** (default) will throw an error if one tries to delete a row (or update the primary key value) referenced.
  - CASCADE** will force the referencing tuples to be deleted (or updated with new primary key value).
  - SET NULL** will force the corresponding values in the referencing tuples to be set to a null value (i.e., unknown).
  - SET DEFAULT** will force the corresponding values in the referencing tuples to be set to a specified default value.



## Delete – Example

- Consider the following foreign key defined on ENROL:

FOREIGN KEY(StudentID) REFERENCES STUDENT(StudentID)  
**ON DELETE NO ACTION**

| ENROL     |          |          |        |            |
|-----------|----------|----------|--------|------------|
| StudentID | CourseNo | Semester | Status | EnrolDate  |
| 456       | COMP1130 | 2016 S1  | active | 25/02/2016 |
| 458       | COMP1130 | 2016 S1  | active | 25/02/2016 |
| 456       | COMP2400 | 2016 S2  | active | 09/03/2016 |

| STUDENT   |       |            |                   |
|-----------|-------|------------|-------------------|
| StudentID | Name  | DoB        | Email             |
| 456       | Tom   | 25/01/1988 | tom@gmail.com     |
| 458       | Peter | 20/02/1991 | peter@hotmail.com |

DELETE FROM STUDENT WHERE StudentID=456;

- The deletion of a student who has enrolled at least one course will throw out an error concerning the foreign key.



## Delete – Example

- Consider the following foreign key defined on ENROL:

FOREIGN KEY(StudentID) REFERENCES STUDENT(StudentID)  
**ON DELETE CASCADE**

| ENROL     |          |          |        |            |
|-----------|----------|----------|--------|------------|
| StudentID | CourseNo | Semester | Status | EnrolDate  |
| 456       | COMP1130 | 2016 S1  | active | 25/02/2016 |
| 458       | COMP1130 | 2016 S1  | active | 25/02/2016 |
| 456       | COMP2400 | 2016 S2  | active | 09/03/2016 |

| STUDENT   |       |            |                   |
|-----------|-------|------------|-------------------|
| StudentID | Name  | DoB        | Email             |
| 456       | Tom   | 25/01/1988 | tom@gmail.com     |
| 458       | Peter | 20/02/1991 | peter@hotmail.com |

DELETE FROM STUDENT WHERE StudentID=456;

- We would have ENROL below after deleting the student 456.

| StudentID | CourseNo | Semester | Status | EnrolDate  |
|-----------|----------|----------|--------|------------|
| 458       | COMP1130 | 2016 S1  | active | 25/02/2016 |



## Week 3

- Quiz feedback
- Insert, delete, update
- Select: simple SQL queries
- Select: advanced SQL queries



## Select Statement

- Select Statement

```
SELECT *  
FROM World  
WHERE "Someone"  
LIKE "%You%"
```



## Select Statement

- The **SELECT** statement has the following basic form:

```
SELECT attribute_list
      FROM table_list
      [WHERE condition]
      [GROUP BY attribute_list [HAVING group_condition]]
      [ORDER BY attribute_list];
```



## Select Statement

| STUDENT   |       |            |                  |
|-----------|-------|------------|------------------|
| StudentID | Name  | DoB        | Email            |
| 456       | Tom   | 25/01/1988 | tom@hotmail.com  |
| 458       | Peter | 23/05/1993 | peter@gmail.com  |
| 459       | Fran  | 11/09/1987 | frankk@gmail.com |

- What is the result for the following Select statement?

```
SELECT * FROM STUDENT WHERE Email like '%@gmail.com';
```

| StudentID | Name  | DoB        | Email            |
|-----------|-------|------------|------------------|
| 458       | Peter | 23/05/1993 | peter@gmail.com  |
| 459       | Fran  | 11/09/1987 | frankk@gmail.com |

```
SELECT StudentID FROM STUDENT WHERE Email like '%@gmail.com';
```

| StudentID |
|-----------|
| 458       |
| 459       |



## Select Statement

| STUDENT   |       |            |                  |
|-----------|-------|------------|------------------|
| StudentID | Name  | DoB        | Email            |
| 456       | Tom   | 25/01/1988 | tom@hotmail.com  |
| 458       | peter | 23/05/1993 | peter@gmail.com  |
| 459       | Fran  | 11/09/1987 | frankk@gmail.com |
| 460       | Peter | 03/09/1992 | Peter@Github.com |

- What is the result for the following Select statement?

```
SELECT * FROM STUDENT WHERE Name = 'Peter';
```

| STUDENT   |       |            |                  |
|-----------|-------|------------|------------------|
| StudentID | Name  | DoB        | Email            |
| 460       | Peter | 03/09/1992 | Peter@Github.com |

```
SELECT * FROM STUDENT WHERE lower(Name) = 'peter';
```

| STUDENT   |       |            |                  |
|-----------|-------|------------|------------------|
| StudentID | Name  | DoB        | Email            |
| 458       | peter | 23/05/1993 | peter@gmail.com  |
| 460       | Peter | 03/09/1992 | Peter@Github.com |



## Select + Group By

- **GROUP BY** *attribute\_list* groups tuples for each value combination in the *attribute\_list*.
- Aggregate functions can be applied to aggregate a group of attribute values into a single value, e.g.,
  - **COUNT** returns the total number of argument values
  - **AVG** returns the average of argument values
  - **MIN** returns the minimum value of the arguments
  - **MAX** returns the maximum value of the arguments
  - **SUM** returns the sum of the argument values
- We can use **HAVING** *condition* to add the condition on the groups.



## Select + Group By – Example

| STUDY     |          |       |
|-----------|----------|-------|
| StudentID | CourseNo | Hours |
| 111       | COMP2400 | 120   |
| 222       | COMP2400 | 115   |
| 333       | STAT2001 | 120   |
| 111       | BUSN2011 | 110   |
| 111       | ECON2102 | 120   |
| 333       | BUSN2011 | 130   |

- What would happen for the following SELECT + Group By StudentID?

```
SELECT ...
FROM STUDY
Group By StudentID;
```



## Select + Group By – Example

| Group     | STUDY     |          |       |
|-----------|-----------|----------|-------|
| StudentID | StudentID | CourseNo | Hours |
| 111       | 111       | COMP2400 | 120   |
|           | 111       | BUSN2011 | 110   |
|           | 111       | ECON2102 | 120   |
| 222       | 222       | COMP2400 | 115   |
| 333       | 333       | STAT2001 | 120   |
|           | 333       | BUSN2011 | 130   |

- What is the result for the following SELECT + Group By StudentID?

```
SELECT StudentID  
FROM STUDY  
Group By StudentID;
```

| StudentID |
|-----------|
| 111       |
| 222       |
| 333       |



## Select + Group By – Example

| Group<br><b>StudentID</b> | STUDY     |          |       |
|---------------------------|-----------|----------|-------|
|                           | StudentID | CourseNo | Hours |
| <b>111</b>                | 111       | COMP2400 | 120   |
|                           | 111       | BUSN2011 | 110   |
|                           | 111       | ECON2102 | 120   |
| <b>222</b>                | 222       | COMP2400 | 115   |
| <b>333</b>                | 333       | STAT2001 | 120   |
|                           | 333       | BUSN2011 | 130   |

- What is the result for the following SELECT + Group By StudentID?

```
SELECT StudentID, COUNT(*)  
FROM STUDY  
Group By StudentID;
```

| StudentID | COUNT |
|-----------|-------|
| 111       | 3     |
| 222       | 1     |
| 333       | 2     |



## Select + Group By – Example

| Group     | STUDY     |          |       |
|-----------|-----------|----------|-------|
| StudentID | StudentID | CourseNo | Hours |
| 111       | 111       | COMP2400 | 120   |
|           | 111       | BUSN2011 | 110   |
|           | 111       | ECON2102 | 120   |
| 222       | 222       | COMP2400 | 115   |
| 333       | 333       | STAT2001 | 120   |
|           | 333       | BUSN2011 | 130   |

- What is the result for the following SELECT + Group By StudentID?

```
SELECT StudentID, MAX(hours)
FROM STUDY
Group By StudentID;
```

| StudentID | MAX |
|-----------|-----|
| 111       | 120 |
| 222       | 115 |
| 333       | 130 |



## Select + Group By – Example

| Group<br><b>StudentID</b> | STUDY     |          |       |
|---------------------------|-----------|----------|-------|
|                           | StudentID | CourseNo | Hours |
| 111                       | 111       | COMP2400 | 120   |
|                           | 111       | BUSN2011 | 110   |
|                           | 111       | ECON2102 | 120   |
| 222                       | 222       | COMP2400 | 115   |
| 333                       | 333       | STAT2001 | 120   |
|                           | 333       | BUSN2011 | 130   |

- What is the result for the following SELECT + Group By StudentID?

```
SELECT StudentID, COUNT(StudentID)
FROM STUDY
Group By StudentID;
```

| StudentID | COUNT |
|-----------|-------|
| 111       | 3     |
| 222       | 1     |
| 333       | 2     |



## Select + Group By – Example

| Group | STUDY     |          |       |
|-------|-----------|----------|-------|
|       | StudentID | CourseNo | Hours |
| 111   | 111       | COMP2400 | 120   |
|       | 111       | BUSN2011 | 110   |
|       | 111       | ECON2102 | 120   |
| 222   | 222       | COMP2400 | 115   |
| 333   | 333       | STAT2001 | 120   |
|       | 333       | BUSN2011 | 130   |

- What is the result for the following SELECT + Group By StudentID?

```
SELECT StudentID, CourseNo  
FROM STUDY  
Group By StudentID;
```

Error Message.



## Select + Group By – Example

| Group     | STUDY     |          |       |
|-----------|-----------|----------|-------|
| StudentID | StudentID | CourseNo | Hours |
| 111       | 111       | COMP2400 | 120   |
|           | 111       | BUSN2011 | 110   |
|           | 111       | ECON2102 | 120   |
| 222       | 222       | COMP2400 | 115   |
| 333       | 333       | STAT2001 | 120   |
|           | 333       | BUSN2011 | 130   |

- What is the result for the following SELECT + Group By StudentID?

```
SELECT *
FROM STUDY
Group By StudentID;
```

Error Message.



## Select + Group By – Example

| Group     | STUDY     |          |       |
|-----------|-----------|----------|-------|
| StudentID | StudentID | CourseNo | Hours |
| 111       | 111       | COMP2400 | 120   |
|           | 111       | BUSN2011 | 110   |
|           | 111       | ECON2102 | 120   |
| 222       | 222       | COMP2400 | 115   |
| 333       | 333       | STAT2001 | 120   |
|           | 333       | BUSN2011 | 130   |

- What is the result for the following SELECT + Group By StudentID?

```
SELECT COUNT(*)  
FROM STUDY  
Group By StudentID;
```

| COUNT |
|-------|
| 3     |
| 1     |
| 2     |



## Select + Group By + Having – Example

| Group    | STUDY     |          |       |
|----------|-----------|----------|-------|
| CourseNo | StudentID | CourseNo | Hours |
| BUSN2011 | 111       | BUSN2011 | 110   |
|          | 333       | BUSN2011 | 130   |
| COMP2400 | 111       | COMP2400 | 120   |
|          | 222       | COMP2400 | 115   |
| ECON2102 | 111       | ECON2102 | 120   |
| STAT2001 | 333       | STAT2001 | 120   |

- What is the result for the following SELECT + Group By + Having?

```
SELECT CourseNo
FROM STUDY
Group By CourseNo
Having MAX(Hours) > 120;
```

| CourseNo |
|----------|
| BUSN2011 |



## Select + Group By + Having – Example

| Group    |     | STUDY     |          |       |
|----------|-----|-----------|----------|-------|
| CourseNo |     | StudentID | CourseNo | Hours |
| BUSN2011 | 111 | BUSN2011  | 110      |       |
|          | 333 | BUSN2011  | 130      |       |
| COMP2400 | 111 | COMP2400  | 120      |       |
|          | 222 | COMP2400  | 115      |       |
| ECON2102 | 111 | ECON2102  | 120      |       |
| STAT2001 | 333 | STAT2001  | 120      |       |

- What is the result for the following SELECT + Group By + Having?

```
SELECT CourseNo
FROM STUDY
Group By CourseNo
Having COUNT(*) > 1;
```

| CourseNo |
|----------|
| BUSN2011 |
| COMP2400 |



## Week 3

- Quiz feedback
- Insert, delete, update
- Select: simple SQL queries
- Select: advanced SQL queries
  - set operations, join operations, subqueries



## A Bunch of Tables

- A Bunch of Tables

**A SQL query walks up to two tables in a restaurant and asks:  
“Mind if I join you?”**



## Set Operations

- SQL incorporates several set operations: **UNION** (set union) and **INTERSECT** (set intersection), and sometimes **EXCEPT** (set difference / minus).
- Set operations result in return of a relation of tuples (no duplicates).
- Set operations apply to relations that have the same attribute types appearing in the same order.



## Set Operations

| STUDY     |          |       |
|-----------|----------|-------|
| StudentID | CourseNo | Hours |
| 111       | COMP2400 | 120   |
| 222       | COMP2400 | 115   |
| 333       | STAT2001 | 120   |
| 111       | BUSN2011 | 110   |
| 111       | ECON2102 | 120   |
| 333       | BUSN2011 | 130   |

- What is the result for the following SQL query?

```
SELECT StudentID FROM STUDY  
WHERE CourseNo='COMP2400'
```

**UNION**

```
SELECT StudentID FROM STUDY  
WHERE CourseNo='ECON2102' ;
```

| StudentID |
|-----------|
| 111       |
| 222       |

**UNION**

| StudentID |
|-----------|
| 111       |



## Set Operations

| STUDY            |                 |       |
|------------------|-----------------|-------|
| <u>StudentID</u> | <u>CourseNo</u> | Hours |
| 111              | COMP2400        | 120   |
| 222              | COMP2400        | 115   |
| 333              | STAT2001        | 120   |
| 111              | BUSN2011        | 110   |
| 111              | ECON2102        | 120   |
| 333              | BUSN2011        | 130   |

- What is the result for the following SQL query?

```
SELECT StudentID FROM STUDY  
WHERE CourseNo='COMP2400'
```

**UNION**

```
SELECT StudentID FROM STUDY  
WHERE CourseNo='ECON2102' ;
```

| StudentID |
|-----------|
| 111       |
| 222       |



## Set Operations

| STUDY     |          |       |
|-----------|----------|-------|
| StudentID | CourseNo | Hours |
| 111       | COMP2400 | 120   |
| 222       | COMP2400 | 115   |
| 333       | STAT2001 | 120   |
| 111       | BUSN2011 | 110   |
| 111       | ECON2102 | 120   |
| 333       | BUSN2011 | 130   |

- What is the result for the following SQL query?

```
SELECT CourseNo FROM STUDY
WHERE StudentID=111
EXCEPT
SELECT CourseNo FROM STUDY
WHERE StudentID=222;
```

| CourseNo |
|----------|
| COMP2400 |
| BUSN2011 |
| ECON2102 |

| CourseNo |
|----------|
| COMP2400 |



## Set Operations

| STUDY     |          |       |
|-----------|----------|-------|
| StudentID | CourseNo | Hours |
| 111       | COMP2400 | 120   |
| 222       | COMP2400 | 115   |
| 333       | STAT2001 | 120   |
| 111       | BUSN2011 | 110   |
| 111       | ECON2102 | 120   |
| 333       | BUSN2011 | 130   |

- What is the result for the following SQL query?

```
SELECT CourseNo FROM STUDY  
WHERE StudentID=111
```

**EXCEPT**

```
SELECT CourseNo FROM STUDY  
WHERE StudentID=222;
```

| CourseNo |
|----------|
| BUSN2011 |
| ECON2102 |



## Set Operations

| STUDY            |                 |       |
|------------------|-----------------|-------|
| <u>StudentID</u> | <u>CourseNo</u> | Hours |
| 111              | COMP2400        | 120   |
| 222              | COMP2400        | 115   |
| 333              | STAT2001        | 120   |
| 111              | BUSN2011        | 110   |
| 111              | ECON2102        | 120   |
| 333              | BUSN2011        | 130   |

- What is the result for the following SQL query?

```
SELECT CourseNo FROM STUDY  
WHERE StudentID=111
```

**EXCEPT**

```
SELECT StudentID FROM STUDY  
WHERE CourseNo='ECON2102' ;
```

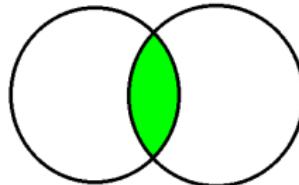
**ERROR MESSAGE**



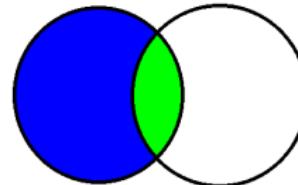
## Join Operations

- When we want to retrieve data from *more than one relations*, we often need to use **join** operations.
- Inner Join:** tuples are included in the result only if there is at least one matching in both relations.
- Left/Right Join:** all tuples of the left/right table are included in the result, even if there are no matches in the relations.

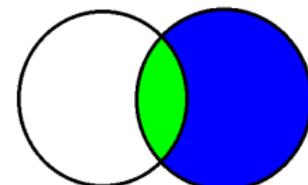
Inner Join



Left Join



Right Join





## Inner Join – Example

| COURSE   |                       |      |
|----------|-----------------------|------|
| No       | Cname                 | Unit |
| COMP2400 | Relational Databases  | 6    |
| BUSN2011 | Management Accounting | 6    |
| ECON2102 | Macroeconomics        | 6    |

| ENROL     |          |          |        |
|-----------|----------|----------|--------|
| StudentID | CourseNo | Semester | Status |
| 111       | BUSN2011 | 2016 S1  | active |
| 222       | COMP2400 | 2016 S1  | active |
| 111       | COMP2400 | 2016 S2  | active |

- What is the result for the following INNER JOIN statement?

SELECT Course.No

FROM COURSE INNER JOIN ENROL ON COURSE.No=ENROL.CourseNo;

| COURSE   |                       |      | ENROL     |          |          |        |
|----------|-----------------------|------|-----------|----------|----------|--------|
| No       | Cname                 | Unit | StudentID | CourseNo | Semester | Status |
| COMP2400 | Relational Databases  | 6    | 222       | COMP2400 | 2016 S1  | active |
| COMP2400 | Relational Databases  | 6    | 111       | COMP2400 | 2016 S2  | active |
| BUSN2011 | Management Accounting | 6    | 111       | BUSN2011 | 2016 S1  | active |



## Inner Join – Example

| COURSE   |                       |      |
|----------|-----------------------|------|
| No       | Cname                 | Unit |
| COMP2400 | Relational Databases  | 6    |
| BUSN2011 | Management Accounting | 6    |
| ECON2102 | Macroeconomics        | 6    |

| ENROL     |          |          |        |
|-----------|----------|----------|--------|
| StudentID | CourseNo | Semester | Status |
| 111       | BUSN2011 | 2016 S1  | active |
| 222       | COMP2400 | 2016 S1  | active |
| 111       | COMP2400 | 2016 S2  | active |

- What is the result for the following INNER JOIN statement?

```
SELECT COURSE.No  
FROM COURSE INNER JOIN ENROL ON COURSE.No=ENROL.CourseNo;
```

| No       |
|----------|
| COMP2400 |
| COMP2400 |
| BUSN2011 |



## Left Join – Example

| COURSE   |                       |      |
|----------|-----------------------|------|
| No       | Cname                 | Unit |
| COMP2400 | Relational Databases  | 6    |
| BUSN2011 | Management Accounting | 6    |
| ECON2102 | Macroeconomics        | 6    |

| ENROL     |          |          |        |
|-----------|----------|----------|--------|
| StudentID | CourseNo | Semester | Status |
| 111       | BUSN2011 | 2016 S1  | active |
| 222       | COMP2400 | 2016 S1  | active |
| 111       | COMP2400 | 2016 S2  | active |

- What is the result for the following LEFT JOIN statement?

SELECT Course.No

FROM COURSE LEFT JOIN ENROL ON COURSE.No=ENROL.CourseNo;

| COURSE   |                       |      | ENROL     |          |          |        |
|----------|-----------------------|------|-----------|----------|----------|--------|
| No       | Cname                 | Unit | StudentID | CourseNo | Semester | Status |
| COMP2400 | Relational Databases  | 6    | 222       | COMP2400 | 2016 S1  | active |
| COMP2400 | Relational Databases  | 6    | 111       | COMP2400 | 2016 S2  | active |
| BUSN2011 | Management Accounting | 6    | 111       | BUSN2011 | 2016 S1  | active |
| ECON2102 | Macroeconomics        | 6    | NULL      | NULL     | NULL     | NULL   |



## Left Join – Example

| COURSE   |                       |      |
|----------|-----------------------|------|
| No       | Cname                 | Unit |
| COMP2400 | Relational Databases  | 6    |
| BUSN2011 | Management Accounting | 6    |
| ECON2102 | Macroeconomics        | 6    |

| ENROL     |          |          |        |
|-----------|----------|----------|--------|
| StudentID | CourseNo | Semester | Status |
| 111       | BUSN2011 | 2016 S1  | active |
| 222       | COMP2400 | 2016 S1  | active |
| 111       | COMP2400 | 2016 S2  | active |

- What is the result for the following LEFT JOIN statement?

```
SELECT Course.No  
FROM COURSE LEFT JOIN ENROL ON COURSE.No=ENROL.CourseNo;
```

| No       |
|----------|
| COMP2400 |
| COMP2400 |
| BUSN2011 |
| ECON2102 |



## Natural Join

- A natural join is considered as one kind of inner join.
- In a natural join, two relations are joined implicitly by comparing all attributes of the same names in both relations.
- A natural join retains all the data of the two tables for only the matched rows, without duplication.



## Natural Join – Example

| COURSE   |                       |      |
|----------|-----------------------|------|
| CourseNo | Cname                 | Unit |
| COMP2400 | Relational Databases  | 6    |
| BUSN2011 | Management Accounting | 6    |
| ECON2102 | Macroeconomics        | 6    |

| ENROL     |          |          |        |
|-----------|----------|----------|--------|
| StudentID | CourseNo | Semester | Status |
| 111       | BUSN2011 | 2016 S1  | active |
| 222       | COMP2400 | 2016 S1  | active |
| 111       | COMP2400 | 2016 S2  | active |

- What is the result for the following NATURAL JOIN statement?

```
SELECT CourseNo  
FROM COURSE NATURAL JOIN ENROL;
```

| CourseNo | COURSE                |      |           | ENROL    |        |  |
|----------|-----------------------|------|-----------|----------|--------|--|
|          | Cname                 | Unit | StudentID | Semester | Status |  |
| COMP2400 | Relational Databases  | 6    | 222       | 2016 S1  | active |  |
| COMP2400 | Relational Databases  | 6    | 111       | 2016 S2  | active |  |
| BUSN2011 | Management Accounting | 6    | 111       | 2016 S1  | active |  |



## Natural Join – Example

| COURSE   |                       |      |
|----------|-----------------------|------|
| CourseNo | Cname                 | Unit |
| COMP2400 | Relational Databases  | 6    |
| BUSN2011 | Management Accounting | 6    |
| ECON2102 | Macroeconomics        | 6    |

| ENROL     |          |          |        |
|-----------|----------|----------|--------|
| StudentID | CourseNo | Semester | Status |
| 111       | BUSN2011 | 2016 S1  | active |
| 222       | COMP2400 | 2016 S1  | active |
| 111       | COMP2400 | 2016 S2  | active |

- What is the result for the following NATURAL JOIN statement?

```
SELECT CourseNo  
FROM COURSE NATURAL JOIN ENROL;
```

| CourseNo |
|----------|
| COMP2400 |
| COMP2400 |
| BUSN2011 |



## Natural Join – Example

| COURSE   |                       |      |
|----------|-----------------------|------|
| No       | Cname                 | Unit |
| COMP2400 | Relational Databases  | 6    |
| BUSN2011 | Management Accounting | 6    |
| ECON2102 | Macroeconomics        | 6    |

| ENROL     |          |          |        |
|-----------|----------|----------|--------|
| StudentID | CourseNo | Semester | Status |
| 111       | BUSN2011 | 2016 S1  | active |
| 222       | COMP2400 | 2016 S1  | active |
| 111       | COMP2400 | 2016 S2  | active |

- What is the result for the following NATURAL JOIN statement?

```
SELECT *
FROM COURSE NATURAL JOIN ENROL;
```

If there are no matching attributes in two tables for NATURAL JOIN,

```
SELECT *
FROM COURSE, ENROL;
```



## Natural Join – Example

| COURSE   |                       |      |
|----------|-----------------------|------|
| CourseNo | Cname                 | Unit |
| COMP2400 | Relational Databases  | 6    |
| BUSN2011 | Management Accounting | 6    |
| ECON2102 | Macroeconomics        | 6    |

| ENROL     |          |          |        |
|-----------|----------|----------|--------|
| StudentID | CourseNo | Semester | Status |
| 111       | BUSN2011 | 2016 S1  | active |
| 222       | COMP2400 | 2016 S1  | active |
| 111       | COMP2400 | 2016 S2  | active |

- What is the result for the following NATURAL JOIN statement?

```
SELECT *
FROM COURSE NATURAL JOIN ENROL ON COURSE.CourseNo=ENROL.CourseNo;
```

**ERROR MESSAGE** because a NATURAL JOIN **implicitly** compares all attributes of the same names in two table.

- Inner Join, Left Join, Right Join, Natural Join, and Cartesian Product, what's their difference?



## Join – More Examples

| STUDENT          |      |     |       |
|------------------|------|-----|-------|
| <u>StudentID</u> | Name | DoB | Email |

| COURSE |       |      |
|--------|-------|------|
| No     | Cname | Unit |

| ENROL            |                 |        |
|------------------|-----------------|--------|
| <u>StudentID</u> | <u>CourseNo</u> | Status |

- List all information of students who have enrolled in a course with CourseNo='X' and the CourseNo of these courses.
  - 1 Use SELECT + FROM (Cartesian Product) + WHERE
  - 2 Use SELECT + FROM (INNER JOIN) + ON
  - 3 Use SELECT + FROM (INNER JOIN) + ON + WHERE
  - 4 Use SELECT + FROM (NATURAL JOIN) + WHERE



## Join – More Examples

| STUDENT   |      |     |       |
|-----------|------|-----|-------|
| StudentID | Name | DoB | Email |

| ENROL     |          |        |
|-----------|----------|--------|
| StudentID | CourseNo | Status |

- List all information of students who have enrolled in a course with CourseNo='X' and the CourseNo of these courses.
- (1) Use SELECT + FROM (Cartesian Product) + WHERE

```
SELECT STUDENT.* , ENROL.CourseNo
FROM STUDENT, ENROL
WHERE (STUDENT.StudentID=ENROL.StudentID)
      AND (ENROL.CourseNo = 'X');
```



## Join – More Examples

| STUDENT   |      |     |       |
|-----------|------|-----|-------|
| StudentID | Name | DoB | Email |

| ENROL     |          |        |
|-----------|----------|--------|
| StudentID | CourseNo | Status |

- List all information of students who have enrolled in a course with CourseNo='X' and the CourseNo of these courses.
- (2) Use SELECT + FROM (INNER JOIN) + ON

```
SELECT STUDENT.* , ENROL.CourseNo
FROM STUDENT INNER JOIN ENROL
ON (STUDENT.StudentID=ENROL.StudentID)
AND (ENROL.CourseNo = 'X');
```



## Join – More Examples

| STUDENT          |      |     |       |
|------------------|------|-----|-------|
| <u>StudentID</u> | Name | DoB | Email |

| ENROL            |                 |        |
|------------------|-----------------|--------|
| <u>StudentID</u> | <u>CourseNo</u> | Status |

- List all information of students who have enrolled in a course with CourseNo='X' and the CourseNo of these courses.
- (3) Use SELECT + FROM (INNER JOIN) + ON + WHERE

```
SELECT STUDENT.* , ENROL.CourseNo
FROM STUDENT INNER JOIN ENROL
ON STUDENT.StudentID=ENROL.StudentID
WHERE ENROL.CourseNo = 'X';
```



## Join – More Examples

| STUDENT   |      |     |       |
|-----------|------|-----|-------|
| StudentID | Name | DoB | Email |

| ENROL     |          |        |
|-----------|----------|--------|
| StudentID | CourseNo | Status |

- List all information of students who have enrolled in a course with CourseNo='X' and the CourseNo of these courses.
- (4) Use SELECT + FROM (NATURAL JOIN) + WHERE

```
SELECT STUDENT.* , ENROL.CourseNo
FROM STUDENT NATURAL JOIN ENROL
WHERE ENROL.CourseNo = 'X';
```



## Subqueries

- **Subqueries** can be viewed as temporary tables (usually in conjunction with aliases and renaming, exist only for the query).
- Subqueries can be specified within the `FROM`-clause.
- Subqueries can also be specified within the `WHERE`-clause, e.g.,
  - `IN subquery` tests if tuple occurs in the temporary table of the subquery.
  - `EXISTS subquery` tests whether the temporary table of the subquery is empty or not.
  - using `ALL`, `SOME` or `ANY` before a subquery makes subqueries usable in comparison formulae (`SOME` and `ANY` are interchangeable).
  - in all these cases the condition involving the subquery can be negated using a preceding `NOT`.



## Subqueries IN – Example 1

| STUDENT          |      |     |       |
|------------------|------|-----|-------|
| <u>StudentID</u> | Name | DoB | Email |

| ENROL            |                 |        |
|------------------|-----------------|--------|
| <u>StudentID</u> | <u>CourseNo</u> | Status |

- List all the information of students who have enrolled in a course with CourseNo='X' and the CourseNo of these courses, we have:

```
SELECT STUDENT.* , ENROL.CourseNo
FROM STUDENT NATURAL JOIN ENROL
WHERE ENROL.CourseNo = 'X';
```



## Subqueries IN – Example 1

- List all the information of students who have enrolled in a course *that has less than 10 students enrolled in it* and the CourseNo of these courses.

- List the CourseNo of the courses in Enrol *that have less than 10 students enrolled*

```
SELECT CourseNo
FROM ENROL
GROUP BY CourseNo
HAVING COUNT(*)<10;
```

- List all the information of students who have enrolled in a course with CourseNo='X' and the CourseNo of these courses

```
SELECT Student.*, Enrol.CourseNo
FROM STUDENT NATURAL JOIN ENROL
WHERE Enrol.CourseNo = 'X';
```



## Subqueries IN – Example 1

- List all the information of students who have enrolled in a course *that has less than 10 students enrolled* and the CourseNo of these courses.

```
SELECT STUDENT.* , CourseNo
  FROM STUDENT NATURAL JOIN ENROL
 WHERE CourseNo IN (SELECT CourseNo
                      FROM ENROL
                     GROUP BY CourseNo
                    HAVING COUNT(*)<10);
```

- Does the above query look confusing?

It is better to distinguish two ENROL tables.



## Subqueries IN – Example 1

- List all the information of students who have enrolled in a course *that has less than 10 students enrolled* and the CourseNo of these courses.
  - Use aliases e1 and e2 for ENROL to distinguish two ENROL tables.

```
SELECT STUDENT.* , e1.CourseNo
FROM STUDENT NATURAL JOIN ENROL e1
WHERE e1.CourseNo IN (SELECT e2.CourseNo
                      FROM ENROL e2
                      GROUP BY e2.CourseNo
                      HAVING COUNT(*) < 10);
```



## Subqueries IN – Example 1

- List all the information of students who have enrolled in a course *that has less than 10 students enrolled* and the CourseNo of these courses.

```
SELECT STUDENT.* , e1.CourseNo
FROM STUDENT NATURAL JOIN ENROL e1
WHERE e1.CourseNo IN (SELECT e2.CourseNo, COUNT(*)
                      FROM ENROL e2
                      GROUP BY e2.CourseNo
                      HAVING COUNT(*) < 10);
```

- Is the above query correct?

No. *IN subquery* tests if tuple occurs in the temporary table of the subquery.

- expression IN (subquery)**

The right-hand side is a parenthesized subquery, which must return exactly one column. (refer to Subquery Expressions in PostgreSQL Documentation)



## Subqueries EXISTS – Example 2

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S1  |
| 222       | COMP2400 | 2016 S1  |
| 111       | COMP2400 | 2016 S2  |

- Count the number of students who have enrolled in at least one course?

```
SELECT COUNT(*)
FROM STUDENT s
WHERE EXISTS (SELECT *
               FROM ENROL e
               WHERE s.StudentID=e.StudentID);
```

1st tuple of STUDENT, EXISTS

| StudentID | CourseNo | Semester |
|-----------|----------|----------|
| 111       | BUSN2011 | 2016 S1  |
| 111       | COMP2400 | 2016 S2  |

2st tuple of STUDENT, EXISTS

| StudentID | CourseNo | Semester |
|-----------|----------|----------|
| 222       | COMP2400 | 2016 S1  |

- The above query (returning 2) is correct!



## Subqueries EXISTS – Example 2

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S1  |
| 222       | COMP2400 | 2016 S1  |
| 111       | COMP2400 | 2016 S2  |

- Count the number of students who have enrolled in at least one course?

```
SELECT COUNT(*)
FROM ENROL e
WHERE EXISTS (SELECT *
               FROM STUDENT s
              WHERE e.StudentID=s.StudentID);
```

1st tuple in ENROL, EXISTS

2nd tuple in ENROL, EXISTS

3rd tuple in ENROL, EXISTS

| StudentID | Name  |
|-----------|-------|
| 111       | Tom   |
| 222       | Emily |
| 111       | Tom   |

- The above query (returning 3 instead of 2) is incorrect!



## Subqueries EXISTS – Example 2

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S1  |
| 222       | COMP2400 | 2016 S1  |
| 111       | COMP2400 | 2016 S2  |

- Count the number of students who have enrolled in at least one course?

```
SELECT COUNT(*)
FROM STUDENT s
WHERE EXISTS (SELECT *
               FROM ENROL e
               WHERE s.StudentID=e.StudentID);
```

```
SELECT COUNT(*)
FROM STUDENT s
WHERE EXISTS (SELECT StudentID
               FROM ENROL e
               WHERE s.StudentID=e.StudentID);
```

- Both queries are correct!** EXISTS subquery tests whether the temporary table of the subquery is empty or not.



## Using Cartesian Product – Example 2

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S1  |
| 222       | COMP2400 | 2016 S1  |
| 111       | COMP2400 | 2016 S2  |

- Count the number of students who have enrolled in at least one course?

```
SELECT COUNT(*)  
FROM STUDENT, ENROL  
WHERE STUDENT.StudentID=ENROL.StudentID;
```

| STUDENT   |       | ENROL     |          |          |
|-----------|-------|-----------|----------|----------|
| StudentID | Name  | StudentID | CourseNo | Semester |
| 111       | Tom   | 111       | BUSN2011 | 2016 S1  |
| 111       | Tom   | 111       | COMP2400 | 2016 S2  |
| 222       | Emily | 222       | COMP2400 | 2016 S1  |

- The above query is incorrect!

We should use COUNT(DISTINCT StudentID) instead of COUNT(\*).



## Using INNER JOIN – Example 2

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S1  |
| 222       | COMP2400 | 2016 S1  |
| 111       | COMP2400 | 2016 S2  |

- Count the number of students who have enrolled in at least one course?

```
SELECT COUNT(*)
FROM STUDENT s INNER JOIN ENROL e
ON s.StudentID=e.StudentID;
```

| s         |       | e         |          |          |
|-----------|-------|-----------|----------|----------|
| StudentID | Name  | StudentID | CourseNo | Semester |
| 111       | Tom   | 111       | BUSN2011 | 2016 S1  |
| 111       | Tom   | 111       | COMP2400 | 2016 S2  |
| 222       | Emily | 222       | COMP2400 | 2016 S1  |

- The above query is incorrect!

We should use COUNT(DISTINCT StudentID) instead of COUNT(\*).



## Using NATURAL JOIN – Example 2

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S1  |
| 222       | COMP2400 | 2016 S1  |
| 111       | COMP2400 | 2016 S2  |

- Count the number of students who have enrolled in at least one course?

```
SELECT COUNT(*)  
FROM STUDENT NATURAL JOIN ENROL;
```

|           | STUDENT | ENROL    |          |
|-----------|---------|----------|----------|
| StudentID | Name    | CourseNo | Semester |
| 111       | Tom     | BUSN2011 | 2016 S1  |
| 111       | Tom     | COMP2400 | 2016 S2  |
| 222       | Emily   | COMP2400 | 2016 S1  |

- The above query is incorrect!

We should use COUNT(DISTINCT StudentID) instead of COUNT(\*).



## A Simple Solution – Example 2

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S1  |
| 222       | COMP2400 | 2016 S1  |
| 111       | COMP2400 | 2016 S2  |

- Count the number of students who have enrolled in at least one course?

```
SELECT COUNT(DISTINCT StudentID)  
FROM ENROL;
```

- The above query is correct!



## Subqueries – Example 3

- List the courses that have the largest number of students enrolled in Semester 2 2016
  - List the CourseNo and the corresponding number of students enrolled for all courses in Semester 2 2016

```
SELECT CourseNo, COUNT(*) AS NoOfStudents
FROM ENROL
WHERE Semester = '2016 S2'
GROUP BY CourseNo;
```

- List **the largest number of students enrolled** in a course in Semester 2 2016

```
SELECT MAX(NoOfStudents)
FROM (SELECT CourseNo, COUNT(*) AS NoOfStudents
      FROM ENROL
      WHERE Semester = '2016 S2'
      GROUP BY CourseNo);
```



## Subqueries – Combination and Aliases – Example 3

- List the courses that have **the largest number of students enrolled** in Semester 2 2016

```
SELECT CourseNo
FROM (SELECT CourseNo, COUNT(*) AS NoOfStudents
      FROM ENROL
      WHERE Semester = '2016 S2'
      GROUP BY CourseNo)
WHERE NoOfStudents =
      (SELECT MAX(NoOfStudents)
       FROM (SELECT CourseNo, COUNT(*) AS NoOfStudents
              FROM ENROL
              WHERE Semester = '2016 S2'
              GROUP BY CourseNo));
```

- ERROR:** Subqueries specifying a derived table must be enclosed in parentheses and must be assigned a table alias name.



## Subqueries – Combination and Aliases – Example 3

- List the courses that have **the largest number of students enrolled** in Semester 2 2016

```
SELECT e.CourseNo
FROM (SELECT e1.CourseNo, COUNT(*) AS NoOfStudents
      FROM ENROL e1
      WHERE e1.Semester = '2016 S2'
      GROUP BY e1.CourseNo) e
WHERE e.NoOfStudents =
      (SELECT MAX(e2.NoOfStudents)
       FROM (SELECT e1.CourseNo, COUNT(*) AS NoOfStudents
              FROM ENROL e1
              WHERE e1.Semester = '2016 S2'
              GROUP BY e1.CourseNo) e2);
```

- Which alias(es) are essential in the above query?  
The aliases e and e2 are essential but e1 is not.



## Subqueries – Use “With” – Example 3

- List the courses that have **the largest number of students enrolled** in Semester 2 2016

Use “WITH” to break down complicated queries into simpler parts.<sup>1</sup>

```
WITH Sem2Students AS
  (SELECT e1.CourseNo, COUNT(*) AS NoOfStudents
   FROM ENROL e1
   WHERE e1.Semester = '2016 S2'
   GROUP BY e1.CourseNo)
SELECT e.CourseNo
  FROM Sem2Students e
 WHERE e.NoOfStudents =
  (SELECT MAX(e2.NoOfStudents)
   FROM Sem2Students e2);
```

- Which alias(es) are essential in the above query?  
None of the aliases e, e1 and e2 are essential.

<sup>1</sup><https://www.postgresql.org/docs/current/static/queries-with.html> 69/78



## Subqueries – Result – Example 3

- List the courses that have **the largest number of students enrolled** in Semester 2 2016

Input:

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S2  |
| 111       | COMP1100 | 2016 S2  |
| 111       | COMP2400 | 2016 S2  |
| 111       | ECON2102 | 2016 S2  |
| 222       | BUSN2011 | 2016 S2  |
| 222       | COMP2400 | 2016 S2  |
| 333       | BUSN2011 | 2016 S2  |
| 333       | COMP2400 | 2016 S2  |
| 333       | ECON2102 | 2016 S2  |

Output:

| CourseNo |
|----------|
| COMP2400 |
| BUSN2011 |



## Subqueries – Example 4

- List all students' IDs and names who are under-enrolled (< 4 courses) in Semester 2 2016.
  - List the students' IDs and the corresponding number of enrolled courses in Semester 2 2016

```
SELECT e.StudentID, COUNT(*) AS NoOfEnrols
FROM ENROL e
WHERE e.Semester = '2016 S2'
GROUP BY e.StudentID;
```



## Subqueries – Example 4

- List all students' IDs and names who are under-enrolled (< 4 courses) in Semester 2 2016.

```
SELECT s.StudentID, s.Name
FROM (SELECT e.StudentID, COUNT(*) AS NoOfEnrols
      FROM ENROL e
      WHERE e.Semester = '2016 S2'
      GROUP BY e.StudentID) ne
INNER JOIN STUDENT s
ON (s.StudentID = ne.StudentID) AND (ne.NoOfEnrols < 4);
```

Use "With"

```
WITH StudEnrols AS (
    SELECT e.StudentID, COUNT(*) AS NoOfEnrols
    FROM ENROL e
    WHERE e.Semester = '2016 S2'
    GROUP BY e.StudentID)
SELECT s.StudentID, s.Name
FROM STUDENT s INNER JOIN StudEnrols ne
ON (s.StudentID = ne.StudentID) AND (ne.NoOfEnrols < 4);
```



## Subqueries - The Query Result - Example 4

- List all students' IDs and names who are under-enrolled (< 4 courses) in Semester 2 2016.

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S2  |
| 111       | COMP1100 | 2016 S2  |
| 111       | COMP2400 | 2016 S2  |
| 111       | ECON2102 | 2016 S2  |
| 222       | BUSN2011 | 2016 S2  |
| 222       | COMP2400 | 2016 S2  |
| 333       | BUSN2011 | 2016 S2  |
| 333       | COMP2400 | 2016 S2  |
| 333       | ECON2102 | 2016 S2  |

Result:

| StudentID | Name  |
|-----------|-------|
| 222       | Emily |
| 333       | John  |

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |



## Subqueries – What About The Following Scenario? – Example 4

- List all students' IDs and names who are under-enrolled (< 4 courses) in Semester 2 2016.

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S2  |
| 111       | COMP1100 | 2016 S2  |
| 111       | COMP2400 | 2016 S2  |
| 111       | ECON2102 | 2016 S2  |
| 222       | BUSN2011 | 2016 S2  |
| 222       | COMP2400 | 2016 S2  |
| 333       | BUSN2011 | 2016 S2  |
| 333       | COMP2400 | 2016 S2  |
| 333       | ECON2102 | 2016 S2  |

Result (still correct?):

| StudentID | Name  |
|-----------|-------|
| 222       | Emily |
| 333       | John  |

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |
| 444       | Ana   |



## Subqueries – Use LEFT/RIGHT JOIN? – Example 4

- List all students' IDs and names who are under-enrolled (< 4 courses) in Semester 2 2016.

```
SELECT s.StudentID, s.Name
FROM (SELECT e.StudentID, COUNT(*) AS NoOfEnrols
      FROM ENROL e
      WHERE e.Semester = '2016 S2'
      GROUP BY e.StudentID) ne RIGHT JOIN STUDENT s
ON (s.StudentID = ne.StudentID) AND (ne.NoOfEnrols < 4);

WITH StudEnrols AS (
    SELECT e.StudentID, COUNT(*) AS NoOfEnrols
    FROM ENROL e
    WHERE e.Semester = '2016 S2'
    GROUP BY e.StudentID)
SELECT s.StudentID, s.Name
FROM STUDENT s LEFT JOIN StudEnrols ne
ON (s.StudentID = ne.StudentID) AND (ne.NoOfEnrols < 4);
```



## Subqueries – Using LEFT/RIGHT JOIN Is Still Incorrect! – Example 4

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S2  |
| 111       | COMP1100 | 2016 S2  |
| 111       | COMP2400 | 2016 S2  |
| 111       | ECON2102 | 2016 S2  |
| 222       | BUSN2011 | 2016 S2  |
| 222       | COMP2400 | 2016 S2  |
| 333       | BUSN2011 | 2016 S2  |
| 333       | COMP2400 | 2016 S2  |
| 333       | ECON2102 | 2016 S2  |

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |
| 444       | Ana   |

Result (still incorrect?):

| StudentID | Name  |
|-----------|-------|
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |
| 444       | Ana   |

- The reason why “111, Tom” is incorrectly included in the final result is due to “Query Processing and Optimisation”, which will be discussed in Week 8. We will revisit this issue in Week 8.



## Subqueries – Use Set Operations – Example 4

- List all students' IDs and names who are under-enrolled (< 4 courses) in Semester 2 2016.

The set of **all students** EXCEPT the set of **students enrolled in at least 4 courses in Semester 2 2016**.

```
SELECT s.StudentID, s.Name
FROM (SELECT StudentID
      FROM STUDENT
      EXCEPT
      SELECT e.StudentID
      FROM ENROL e
      WHERE e.Semester = '2016 S2'
      GROUP BY e.StudentID
      HAVING COUNT(*) > 3) e4 INNER JOIN Student s
ON (e4.StudentID = s.StudentID);
```



## Subqueries – Using Set Operations Works – Example 4

- List all students' IDs and names who are under-enrolled (< 4 courses) in Semester 2 2016.

| ENROL     |          |          |
|-----------|----------|----------|
| StudentID | CourseNo | Semester |
| 111       | BUSN2011 | 2016 S2  |
| 111       | COMP1100 | 2016 S2  |
| 111       | COMP2400 | 2016 S2  |
| 111       | ECON2102 | 2016 S2  |
| 222       | BUSN2011 | 2016 S2  |
| 222       | COMP2400 | 2016 S2  |
| 333       | BUSN2011 | 2016 S2  |
| 333       | COMP2400 | 2016 S2  |
| 333       | ECON2102 | 2016 S2  |

Result:

| StudentID | Name  |
|-----------|-------|
| 222       | Emily |
| 333       | John  |
| 444       | Ana   |

| STUDENT   |       |
|-----------|-------|
| StudentID | Name  |
| 111       | Tom   |
| 222       | Emily |
| 333       | John  |
| 444       | Ana   |